



Direct Method: Sequential Covering

1. Start from an empty rule
2. Grow a rule using the Learn-One-Rule function
3. Remove training records covered by the rule
4. Repeat Step (2) and (3) until stopping criterion is met

Example of Sequential Covering




(i) Original Data




(ii) Step 1

Example of Sequential Covering...




(iii) Step 2



(iv) Step 3


Instance Elimination

- Why do we need to eliminate instances?
 - Otherwise, the next rule is identical to previous rule
- Why do we remove positive instances?
 - Ensure that the next rule is different
- Why do we remove negative instances?
 - Prevent underestimating accuracy of rule
 - Compare rules R2 and R3 in the diagram




Rule Growing

□ Two common strategies



(a) General-to-specific



(b) Specific-to-general

Rule Evaluation

□ Foil's Information Gain

FOIL: First Order Inductive Learner – an early rule-based learning algorithm

- $R0: \{\} \Rightarrow \text{class}$ (initial rule)
- $R1: \{A\} \Rightarrow \text{class}$ (rule after adding conjunct)
- $\text{Gain}(R0, R1) = t [\log (p1/(p1+n1)) - \log (p0/(p0 + n0))]$
- where t : number of positive instances covered by both $R0$ and $R1$
 $p0$: number of positive instances covered by $R0$
 $n0$: number of negative instances covered by $R0$
 $p1$: number of positive instances covered by $R1$
 $n1$: number of negative instances covered by $R1$

Direct Method: RIPPER

- For 2-class problem, choose one of the classes as positive class, and the other as negative class
 - Learn rules for positive class
 - Negative class will be default class
- For multi-class problem
 - Order the classes according to increasing class prevalence (fraction of instances that belong to a particular class)
 - Learn the rule set for smallest class first, treat the rest as negative class
 - Repeat with next smallest class as positive class

Direct Method: RIPPER

- Growing a rule:
 - Start from empty rule
 - Add conjuncts as long as they improve FOIL's information gain
 - Stop when rule no longer covers negative examples
 - Prune the rule immediately using incremental reduced error pruning
 - Measure for pruning: $v = (p-n)/(p+n)$
 - ◆ p : number of positive examples covered by the rule in the validation set
 - ◆ n : number of negative examples covered by the rule in the validation set
 - Pruning method: delete any final sequence of conditions that maximizes v

Direct Method: RIPPER

□ Building a Rule Set:

- Use sequential covering algorithm
 - ◆ Finds the best rule that covers the current set of positive examples
 - ◆ Eliminate both positive and negative examples covered by the rule
- Each time a rule is added to the rule set, compute the new description length
 - ◆ Stop adding new rules when the new description length is d bits longer than the smallest description length obtained so far

Direct Method: RIPPER

- Optimize the rule set:
 - For each rule r in the rule set R
 - ◆ Consider 2 alternative rules:
 - Replacement rule (r^*): grow new rule from scratch
 - Revised rule(r'): add conjuncts to extend the rule r
 - ◆ Compare the rule set for r against the rule set for r^* and r'
 - ◆ Choose rule set that minimizes MDL principle
 - Repeat rule generation and rule optimization for the remaining positive examples